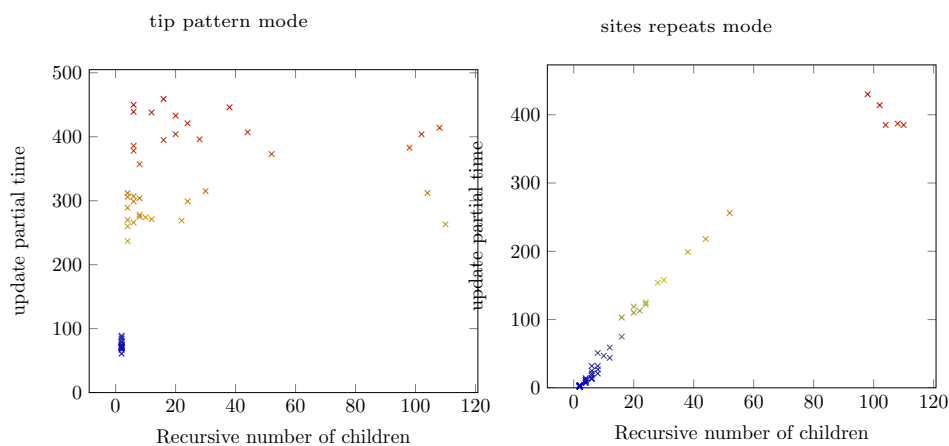


Libpll per node benchmarks


December 7, 2016

Each point on the plots corresponds to one inner node. We first compute the recursive number of children under this node (x axis). Then we run update partials on this node (1000 times) and plot the elapsed time (y axis). We do that for each dataset and with/without sites repeats.

seq 59



seq 128



seq 404

